

ECE 473/573
Cloud Computing and Cloud Native Systems
Lecture 11 Database Systems

Professor Jia Wang
Department of Electrical and Computer Engineering
Illinois Institute of Technology

September 22, 2025

Outline

Cloud Storage

Relational Database

Relational Algebra and SQL

Reading Assignment

- ▶ This lecture: Database systems
- ▶ Next two lectures: Distributed database systems
 - ▶ Cassandra - A Decentralized Structured Storage System
<https://www.cs.cornell.edu/projects/ladis2009/papers/lakshman-ladis2009.pdf>
 - ▶ Spanner: Google's Globally-Distributed Database
http://static.googleusercontent.com/external_content/untrusted_dlcp/research.google.com/en/archive/spanner-osdi2012.pdf

Outline

Cloud Storage

Relational Database

Relational Algebra and SQL

Cloud Storage

- ▶ A fundamental component of cloud computing.
 - ▶ Persist state of microservices and applications.
 - ▶ Store intermediate data to facilitate communication and fault resilience.
- ▶ Metrics
 - ▶ Size
 - ▶ Performance: throughput and latency
 - ▶ Availability and reliability
 - ▶ Leverage scalability to improve all of them.
- ▶ Different types of cloud storage may have different trade-offs.
 - ▶ Block storage and file systems
 - ▶ Object storage
 - ▶ Database systems

Block Storage and File Systems

- ▶ Block storage provides byte blocks of fixed size that can be accessed randomly.
 - ▶ E.g. hard drives and solid-state drives.
 - ▶ Available locally or through a dedicated network (SAN) for high throughput and low latency.
- ▶ File systems built on top of block storage provide support to
 - ▶ Organize data as files and directories
 - ▶ Share files over network
 - ▶ Checksum, versioning, and redundancy
 - ▶ Security features like permission and encryption
- ▶ Not scalable
 - ▶ Strong tie to the underlying hardware for performance
 - ▶ Exclusive access is required to update a block.

Object Storage

- ▶ Manage data as objects that must be modified as a whole.
 - ▶ Accessed via networked services.
 - ▶ Use a key as identifier instead of a name.
 - ▶ Need other mechanisms to support hierarchy.
- ▶ Highly scalable
 - ▶ Able to utilize physical storage from many servers via networked services.
 - ▶ Many objects are not modified after creation – easy to maintain multiple copies of the same object.
- ▶ Optimize for different access patterns
 - ▶ Backups that are mostly write-once without read.
 - ▶ Intermediate data that require only sequential access.
 - ▶ Media files that are mostly read-only but need to be read frequently from all over the world.

Database Systems

- ▶ Provide rich accesses to highly structured data beyond read/write.
- ▶ Relational (SQL) database
 - ▶ Very strong guarantee on data consistency – a must to manage data that need to be consistent like payments.
 - ▶ Mature and well-understood.
 - ▶ Not scalable – need to maintain a lot of internal states.
- ▶ NoSQL databases
 - ▶ High scalability by giving up some part of the consistency guarantees of SQL databases.
 - ▶ Different NoSQL databases may explore different trade-offs to favorite different applications, making it tricky to pick up the right one to meet the requirement.

Outline

Cloud Storage

Relational Database

Relational Algebra and SQL

Relational Database

- ▶ A classical approach for data management.
 - ▶ Restrict functionality to what can be expressed in relational algebra, usually captured by the SQL language.
 - ▶ Provide ACID guarantee on database operations including data persistency and concurrent access.
- ▶ Usually run as a stand-alone service that clients can access locally or remotely.
 - ▶ Via management tools, or
 - ▶ Via APIs that are available from most programming languages.

ACID Guarantee

- ▶ Database updates are grouped into transactions to support application logic.
 - ▶ E.g. if Alice need to transfer \$100 to Bob, the transaction need to deduct \$100 from Alice's account and add \$100 to Bob's account.
- ▶ **A**tomicity: either the transaction succeeds or fails as a whole.
 - ▶ It is not allowed to deduct \$100 from Alice's account while not changing Bob's account.
- ▶ **C**onsistency: database remains valid after transactions are executed.
 - ▶ Transactions are committed if succeed. Later transactions will see the changes.
 - ▶ Failed transactions should not change the database.
 - ▶ Transactions, if committed, should execute correctly, e.g. it is not allowed to deduct \$100 from Alice's account while adding \$50 to Bob's account, and not allowed for Alice to have a negative balance.

ACID Guarantee (Cont.)

- ▶ **I**solaton: transactions are executed as if sequentially.
 - ▶ Actual implementations may execute transactions concurrently to achieve better performance.
 - ▶ However, the outcome should be the same as if the transactions are executed one after another – note that the order is not specified.
 - ▶ E.g. if we assume Alice initially has \$0 in her account and that at the same time Alice transfers \$100 to Bob, Carol transfers \$200 to Alice, then both are possible that the transaction from Alice to Bob succeeds or fails.
- ▶ **D**urability: committed transactions survive system failures.
 - ▶ Usually by storing data on a drive.
 - ▶ To the extent that the drive won't fail.
- ▶ It is quite challenge to achieve ACID at the same time.
 - ▶ E.g. what if there is a power outage when the database is about to commit one transaction by writing data to the disks?

Data Models in Relational Database

- ▶ Data are organized into tables or relations.
- ▶ Each table consists of many rows or tuples of data.
- ▶ Each row consists of many columns or attributes or fields
 - ▶ Rows in the same table should have the same columns.
- ▶ Each row should have a special column called the key or the primary key that is unique among the rows in the same table.
 - ▶ Allow one to quickly locate the row given its key.
 - ▶ Additionally to support a range query of keys.
- ▶ Each column of a row is usually of an elementary data type.
 - ▶ That can be compared and operated on.
 - ▶ Opaque binary blobs are also supported by many database systems to store data like images.

Outline

Cloud Storage

Relational Database

Relational Algebra and SQL

SQL Query

```
SELECT users.id, SUM(orders.total) total_spending,  
FROM users JOIN orders ON (users.id=orders.buyer_id)  
WHERE orders.year=2023  
GROUP BY users.id  
ORDER BY total_spending DESC;
```

- ▶ SQL queries start with the **SELECT** clause.
- ▶ Each query will return rows of data.
 - ▶ Each row may contain data from multiple tables.
 - ▶ Columns are specified in the **SELECT** clause.
 - ▶ E.g. two columns **users.id** and **total_spending** are generated here.

Data Source

```
SELECT ...  
FROM users JOIN orders ON (users.id=orders.buyer_id)  
...
```

- ▶ The **FROM** clause specifies data to query from.
- ▶ You may query data from a single table, or
- ▶ From multiple tables by joining them together.
 - ▶ So that relevant data can be retrieved from multiple tables at the same time.

Join

```
SELECT ...  
FROM users JOIN orders ON (users.id=orders.buyer_id)  
...
```

- ▶ There are many kinds of **JOINS**: one method to understand all of them is to consider **JOIN** as a two-step process.
- ▶ Step 1: form a new table by taking the Cartesian product of the tables.
 - ▶ If **users** has N rows and **orders** has M rows, the new table will have NM rows, each consists of a row from **users** and a row from **orders**.
- ▶ Step 2: remove rows from the new table following certain criteria as defined by different **JOINS**.
 - ▶ For the above example, we remove the rows where **users.id** and **orders.buyer_id** are different.
 - ▶ The new table lists buyers and their orders together.
- ▶ Actual implementations may eliminate the need to calculate the Cartesian product depending if the criteria involves keys.

Filtering

```
SELECT ...  
FROM ...  
WHERE orders.year=2023  
...
```

- ▶ The **WHERE** clause filters rows by a given condition.
 - ▶ So that a portion of the whole table may be retrieved.
 - ▶ E.g. for this query we only care about **orders** placed in 2023.

Grouping and Aggregation

```
SELECT users.id, SUM(orders.total) total_spending,  
FROM ...  
WHERE ...  
GROUP BY users.id  
...
```

- ▶ Rows in the joined new table may be further grouped via **GROUP BY** clause.
 - ▶ E.g. to group all rows belonging to the same buyer together.
- ▶ As SQL only operates on rows but not groups of rows, rows from each group must be aggregated into a new row.
 - ▶ Via aggregate functions like **SUM** to calculate the total spending of each buyer.

Output Ordering

```
SELECT users.id, SUM(orders.total) total_spending,  
FROM users JOIN orders ON (users.id=orders.buyer_id)  
WHERE orders.year=2023  
GROUP BY users.id  
ORDER BY total_spending DESC;
```

- ▶ Finally, the output rows may be sorted via **ORDER BY**.
 - ▶ Either ascending (**ASC**) or descending (**DESC**).
 - ▶ So that we can find who spends the most for 2023.

Other SQL Statements

- ▶ There are other SQL statements to create, update, and delete rows from tables and to manage tables as well.
- ▶ Check <https://www.w3schools.com/sql/default.asp> and run examples there to learn SQL.