

ECE 473/573  
Cloud Computing and Cloud Native Systems  
Lecture 13 Distributed Database Systems II

Professor Jia Wang  
Department of Electrical and Computer Engineering  
Illinois Institute of Technology

September 29, 2025

# Outline

Spanner

# Reading Assignment

- ▶ This lecture: Spanner: Google's Globally-Distributed Database [http://static.googleusercontent.com/external\\_content/untrusted\\_dlcp/research.google.com/en//archive/spanner-osdi2012.pdf](http://static.googleusercontent.com/external_content/untrusted_dlcp/research.google.com/en//archive/spanner-osdi2012.pdf)
- ▶ Next lecture: Introduction to Cloud Security

# Outline

Spanner

- ▶ A distributed multi-version database.
  - ▶ Semi-relational with SQL support.
  - ▶ Transactions with ACID guarantee.
  - ▶ Globally-distributed and horizontally scalable.
- ▶ Sharding automatically help to balance loads as data grow and as servers join and leave the cluster.
- ▶ Replication for global availability and geographic locality.
  - ▶ Configurable by applications.
  - ▶ Location of replicas: read latency, write latency.
  - ▶ Number of replicas: read performance, durability and availability.

# Spanner Deployment

- ▶ A Spanner deployment is called a universe.
  - ▶ E.g. a universe for testing, and another for production.
- ▶ Physical servers are managed as zones.
  - ▶ Each zone is the unit of administration and physical isolation.
  - ▶ E.g. a datacenter may contain multiple zones, one for each application that need to be isolated.
- ▶ Each zone has one zonemaster and a number (hundreds to thousands) of spanservers.
  - ▶ Zonemaster assigns data to spanservers.
  - ▶ Spanservers serve data to clients.

# Tablets, Directories, and Placement

- ▶ Each spanserver manages many tablets.
- ▶ Each tablet maintains many versioned key-value pairs.
  - ▶ I.e., past updates to a key-value pair are recorded.
  - ▶ Tablets are replicated across many spanservers.
- ▶ Key-value pairs within a tablet are grouped into directories.
  - ▶ Keys in each directory share a common prefix.
- ▶ In other words, the common prefix determines where the key-value pairs are stored and replicated.
  - ▶ A directory is the smallest unit whose placement, i.e. geographic replication properties like 5 replicas in US, can be configured by applications.

# Data Model

- ▶ Similar to relational databases.
- ▶ An application can create multiple databases in a universe and each database consists of multiple schematized tables.
- ▶ Each table consists of rows and each row has a predefined list of columns, some as the key and the rest as the value.
  - ▶ Each row corresponds to a key-value pair in a tablet so its update history is recorded.
- ▶ Key columns are ordered and part of them are use for the common prefix defining the directory this row belongs to.



## Data Model (Cont.)

- ▶ Unlike Cassandra, Spanner supports SQL features like joins.
- ▶ Support transactions across rows in a distributed manner.
- ▶ Provide consistency and partition tolerance, while let applications handles availability issues.
- ▶ Indeed, the CAP theorem says it is not possible to have 100% availability with consistency and partition tolerance, but in practice we don't always need 100% availability.

# Consistency across Replicas of the Same Row

- ▶ Consensus: if multiple writes to the same row arrive at different servers, which one will succeed?
- ▶ For Cassandra, eventually consistent requires all replicas of the same row to be the same eventually when there is no more writes.
  - ▶ If writes are not acknowledged from all replicas, then there is no guarantee reads from the replica not acknowledged will return the same as reading other replicas – no consensus at all.
- ▶ For Spanner, a consensus protocol ensures replicas of the same tablet across multiple spanservers record the same history.
  - ▶ Reading any replica will give the same history of writes – only some replicas have more recent history than others.
  - ▶ We will introduce the consensus protocol Paxos toward the end of the semester.

# Cross-Row Transactions

- ▶ However, the consensus protocol does not guarantee anything for writes to different rows not in the same tablet.
- ▶ Recall our social network example.
  - ▶ TABLE Friends stores rows for friendship relation.
  - ▶ TABLE Posts stores rows for posts.
- ▶ If a user A removes a friend X and then creates a post P, then A does not want X to read P.
- ▶ Three transactions are of interests for this scenario.
  - ▶ A0: remove X from A's friend list and A from X's friend list.
  - ▶ A1: add P to A's posts.
  - ▶ X0: read friend list of X, then list posts for each friend of X.
- ▶ A0 and A1 need to write to rows in different tablets and replicas and X0 need to read them.
  - ▶ The replicas containing A's friend list.
  - ▶ The replicas containing X's friend list.
  - ▶ The replicas containing A's posts.

## Cross-Row Transactions (cont.)

- ▶ What if X0 reads a more recent replica with A's posts than a replica with X's friend list?
  - ▶ Output of X0 will include P which it should not.
  - ▶ As if A1 completes before A0.
- ▶ No, one cannot wait for all replicas to have the most recent data before executing X0.
  - ▶ There may be other transactions updating the replicas constantly.
  - ▶ Those transactions have to run concurrently, and cannot be blocked for availability and performance reasons.
- ▶ How does traditional relational database solve this problem with ACID guarantees?
- ▶ What prevents distributed databases to do the same?

# ACID Guarantees

- ▶ Traditional relational databases provide ACID guarantees.
  - ▶ We can understand the overall effects of these transactions by inspecting all possible orderings assuming they execute and complete one after another.
- ▶ Six possible orderings of A0, A1, and X0
  - ▶ Three orderings have A0 completes before A1
  - ▶ The other three have A1 completes before A0
- ▶ For the correctness of transaction execution, we would expect the three with A0 before A1.
  - ▶ X0, A0, A1: X only sees post for A before A removes X
  - ▶ A0, X0, A1: X don't see any post from A
  - ▶ A0, A1, X0: A posts P but X don't see any post from A
- ▶ What prevents A1 to complete before A0?
  - ▶ Time causality on the single server: since user A wait for A0 to complete before starting A1, a local clock on that single server ensures A1 to start after A0 completes.

# External Consistency

- ▶ In a distributed database when A0 and A1 write to rows on different servers, these servers have different local clocks.
  - ▶ X1 may see A1 completes before A0 using their local timestamps.
- ▶ External consistency: still, from the viewpoint of A's local clock, A0 does complete before A1 starts.
  - ▶ But how can such timing information be incorporated, which is external to the database system, into the transactions?
- ▶ Will it help if all local clocks synchronize with a global clock?
- ▶ Does such a global clock exist at all?

# Version Data and Global Clock

- ▶ Since Spanner keeps versioned data, if the versioned data are stamped with a global clock, here is a possible solution.
  - ▶ X0's query into A's posts depending on a query on X's friend list. Therefore, it should not use any data from A's post more recent than from X's friend list.
  - ▶ X0 can be thought to happen sometime back in the history and correctness is achieved!
- ▶ For multiple transactions reading X's friend list,
  - ▶ Reading different replicas will result in different times back in the history those transactions thought to happen.
  - ▶ The consensus on the history among all replicas ensures their outcomes to follow external consistency.
- ▶ Can we maintain a global clock that multiple servers distributed to different locations can synchronize with?
  - ▶ But special relativity says there is no such global clock.

- ▶ GPS and atomic clocks can provide accurate time for local clocks and can compensate for each other as they have different failure modes.
- ▶ With an algorithm to synchronize time between local clocks, we can have the illusion of a global clock.
  - ▶ Each local clock has a time uncertainty with respect to the global clock that can be measured.
  - ▶ No violation of special relativity since uncertainty will increase as distances increase.
- ▶ Each transaction will use the local clock to stamp its writes.
  - ▶ To ensure that the timestamps from transactions to follow their commit order, transactions will need to wait twice of the uncertainty bound.



# TrueTime Example

- ▶ Global clock uncertainty: 500ms
  - ▶ Local clocks on servers are less than 1s away from each other.
  - ▶ Servers have no other knowledge of local clocks of each other.
- ▶ Consider two servers
  - ▶ XF: the one containing X's friend list.
  - ▶ AP: the one containing A's posts.
- ▶ First A is removed from X's friend list
  - ▶ Stamped with local time of XF: 8:00:00.000
  - ▶ Local time of AP: 7:59:59.001
- ▶ Then P is added to A's posts.
  - ▶ Local time of AP should be at least 7:59:59.001
  - ▶ It is incorrect to stamp the event with 7:59:59.001.
  - ▶ Wait 2x500ms and stamp the event with 8:00:00.001.
- ▶ All queries now see P is added after A is removed.
- ▶ What if local time of AP is 8:00:00.999 when local time of XF is 8:00:00.000?