

ECE 473/573  
Cloud Computing and Cloud Native Systems  
Lecture 21 Message Queues

Professor Jia Wang  
Department of Electrical and Computer Engineering  
Illinois Institute of Technology

October 29, 2025

# Outline

Message Queues

Kafka

# Reading Assignment

- ▶ This lecture: Apache Kafka  
<https://kafka.apache.org/documentation/#design>
- ▶ Next Lecture: 9

# Outline

Message Queues

Kafka

# Message Queues

- ▶ A middleware to enable message communication between senders and receivers, e.g. a message broker.
  - ▶ Reduce coupling by removing the immediate connections between message senders and receivers.
  - ▶ Serve as a buffer to reduce impact of performance difference between senders and receivers when there is a burst of load.
- ▶ How messages are distributed among receivers for a queue?
  - ▶ Producer-Consumer: senders are producers generating jobs as messages, receivers are consumers taking jobs out of the queue to work on them – no two consumers work on the same job.
  - ▶ Publisher-Subscriber (Pub-Sub): senders publish messages and all receivers as subscribers receive all messages.

# Persistence

- ▶ What if queues fails?
- ▶ Persistence is required for producer-consumer queues.
  - ▶ Otherwise jobs may be lost.
- ▶ Persistence for Pub-Sub queues
  - ▶ With persistence, a subscriber can subscribe at any time to receive all past and future messages.
  - ▶ Without persistence, a subscriber only receives future messages after it subscribes but not past ones.
  - ▶ However, it takes resources to support persistence so one need to make a choice depending on application requirements.

# Topic Management

- ▶ Queues are usually identified by topics.
  - ▶ Usually a meaningful string providing hints on what messages inside are all about.
  - ▶ More frequently used for Pub-Sub queues.
- ▶ Each publisher or sender, sends (topic, message) to the message broker.
  - ▶ So the message broker knows to which queue the message should go.
- ▶ Each subscriber or receiver, when establishing connection with the message broker, specifies what topics it is interested into.
  - ▶ Then the message broker will only send (topic, message) with matching topics to this subscriber.

# Ideal Queue Behavior

- ▶ FIFO (first-in-first-out) ordering.
  - ▶ Messages are delivered to the receivers in the order they are sent by the senders.
- ▶ Exactly-once delivery.
  - ▶ Messages sent by senders are delivered to receivers exactly once – no lost messages and no repetition.
  - ▶ For producer-consumer queues, a producer generates a message and exactly one consumer consumes it exactly once.
  - ▶ For publisher-subscriber queues, a publisher publishes a message and every subscriber receives it exactly once.
- ▶ Can we achieve these with networked services?
  - ▶ With a single message broker.
  - ▶ With multiple message brokers for horizontal scalability.



# Communication Challenges for Ordering

- ▶ Consider when there is a single message broker.
  - ▶ The broker should be able to store and then send out messages in the order of their arrivals.
  - ▶ The FIFO ordering makes it possible to optimize persistence for high sequential read and write throughput.
- ▶ It takes time for messages to arrive from senders to the message broker and from the message broker to receivers.
  - ▶ For best performance, messages may take different network paths so may arrive out-of-order.
  - ▶ Messages from different senders on the same topic can only be ordered when they arrive at the message broker.
- ▶ Why can't the message broker reorder messages differently than the arrival order, e.g. using timestamps as keys?
  - ▶ This functionality should be provided through a key-value database that is more complicated and much slower.

# Communication Challenges for Exactly-Once Delivery

- ▶ Network communications are unreliable.
  - ▶ Protocols like TCP and HTTP guarantee delivery only when there is no failure.
- ▶ Delivery guarantees considering failures
  - ▶ At-most once (best effort): messages are sent once, don't use any acknowledgement.
  - ▶ At-least once: resend messages until acknowledgements are received.
- ▶ Apparently, exactly-once delivery can be achieved by using sequence numbers with at-least once delivery.
  - ▶ Message brokers number messages as they arrive.
  - ▶ Subscribers and consumers utilize these numbers to acknowledge and reorder messages.
- ▶ However, maintaining sequence numbers as messages are generated by publishers and producers is not trivial.
  - ▶ Resending unacknowledged messages further complicates the issue as it leads to additional out-of-order arrivals.

# Scalability Challenges

- ▶ Consider horizontal scalability where multiple message brokers are running on multiple servers.
- ▶ Each message broker can handle a number of topics.
  - ▶ Similar to the idea of sharding and function partitioning.
- ▶ Each topic can be replicated to multiple message brokers.
  - ▶ Which then serve huge number of subscribers.
- ▶ However,
  - ▶ Scaling with multiple publishers and producers sending messages to the same topic is difficult.
  - ▶ Scaling for consumers is difficult as replicas need to have consensus on which consumer should process which message.

# Outline

Message Queues

Kafka

- ▶ An open-source distributed event streaming platform.
  - ▶ Developed in LinkedIn, open-sourced in 2011
- ▶ Features
  - ▶ Low-latency message delivery for high volume event streams, e.g. real-time log aggregation and offline data loading.
  - ▶ Support a computational model for real-time analytics by consuming and producing event streams.
  - ▶ Fault tolerance.

# Architecture

- ▶ Events as messages are organized and stored in topics.
- ▶ A combination of producer-consumer and publisher-subscriber message queues for scalability.
  - ▶ Kafka producers publish (write) events to topics.
  - ▶ Kafka consumers subscribe to topics and read events within.
  - ▶ Events in a topic are partitioned – a group of consumers can read these events in parallel, each for a different partition.
  - ▶ Multiple groups of consumers can still read events in a topic as many times as desired.
- ▶ Kafka brokers store partitions for topics.
  - ▶ Sharding: partitions of a single topic are distributed to multiple brokers for better write performance.
  - ▶ Replication: a single partition is replicated across multiple brokers for better read performance, availability, and durability.

# Persistence and Performance

- ▶ Major factor to drive Kafka design decisions.
- ▶ Rely heavily on filesystem for storing and caching messages.
  - ▶ Sequential write and read throughput are high enough to saturate network communications.
  - ▶ OS automatically makes efficient use of large amount of memory when caching disk data for sequential accesses.
  - ▶ On the contrary, most languages cannot use memory as efficiently due to object overhead and garbage collections.
- ▶ Disk space is virtually unlimited so messages can be kept for a long time before being deleted.
- ▶ To guarantee high performance, only rely on simple and sequential accesses to files.
  - ▶ Store messages by appending to files.
  - ▶ Serve messages by reading sequentially.
  - ▶ Not to support any kinds of indices that would need random accesses – use databases instead.

# Message Delivery Semantics

- ▶ Assume brokers to work perfectly for now.
- ▶ No guarantee among multiple producers within Kafka.
- ▶ No ordering among multiple partitions, even if they are from the same topic.
- ▶ Exactly-once processing is supported via Kafka transactions.
  - ▶ A single transaction reads from and writes to multiple partitions, possibly from different topics.
- ▶ There are options for weaker guarantees for other use cases.
  - ▶ A single producer with a single partition.
  - ▶ A group of consumers with a topic.



# Message Delivery for Producers

- ▶ A single producer with a single partition.
- ▶ At-most once: broker doesn't acknowledge
  - ▶ Messages may arrive out of order.
- ▶ At-least once: producer resends messages until acknowledged
  - ▶ Broker may store a message twice.
  - ▶ If a previous message is not acknowledged yet, the next message may arrive out of order.
- ▶ Idempotent delivery: producer adds sequence numbers
  - ▶ Base on at-least once delivery
  - ▶ Broker remove duplicates and acknowledges messages in-order.
  - ▶ May affect performance as out-of-order arrivals are not acknowledged and need resend.
  - ▶ Only for the lifetime of the producer, no guarantee if it restarts.

# Message Processing for Consumers

- ▶ A group of consumers with a topic.
- ▶ Each consumer reads one partition of the topic.
- ▶ Each consumer saves its position within its partition.
  - ▶ The position indicates where to read from if the consumer restarts after a failure.
  - ▶ However, there are two choices as the consumer need to process the message.
- ▶ Process-then-save: at-least once
  - ▶ If the consumer fails after processing the message but before saving the position, then when it restarts, it will process the message again.
  - ▶ Make sure the processing is idempotent to avoid any issues.
- ▶ Save-then-process: at-most once
  - ▶ If the consumer fails after saving the position but before processing the message, then when it restarts, it will skip the message.

# Replication

- ▶ Unit of replication is a partition.
- ▶ Consensus determines a leader replica for each partition.
  - ▶ Producers write to leader directly.
  - ▶ Other replicas (followers) replicate from the leader.
  - ▶ Consensus and all states are managed in ZooKeeper.
  - ▶ Not P for CAP theorem: no partition tolerance.
- ▶ In-sync replicas (ISRs)
  - ▶ Replicas that are not too far behind the leader.
  - ▶ Messages available from all ISRs are considered committed.
  - ▶ Committed messages are less likely to be lost if the leader fails.
  - ▶ Consumers only consume committed messages.
  - ▶ Producers can choose to receive acknowledgement when the message reaches the leader or when it is committed.

# Summary

- ▶ Use message queues to decouple message senders and receivers.
- ▶ Make a choice between distributed message queues and distributed database systems by considering their performance differences and your application use cases.